

CCDA : Credit Card Data Analysis

1.Database Design



01.ccda_final_design.pdf

2. Create AWS RDS instance of Postgre (PostgreSQL 12.5-R1) type

Note : Make DB public , remember Port,Username,Password and Endpoint

Connectivity & security

Monitoring

Logs & events

Configuration

Maintenance & backups

Tags

Connectivity & security

Endpoint & port

Endpoint

database-1.cjpk2s29ue0.ap-south-1.rds.amazonaws.com

Port

5432

Networking

Availability zone

ap-south-1b

VPC

vpc-0b7c4f63

Subnet group

default-vpc-0b7c4f63

Subnets

subnet-65ba3429

subnet-1fbb6b64

subnet-75ca931d

Security

VPC security groups

default (sg-d8f223b7)

(active)

Public accessibility

Yes

Certificate authority

rds-ca-2019

Certificate authority date

August 22, 2024, 10:38 (UTC+10:38)

3. Connect That Created Postgre instance using workbench

Driver

PostgreSQL (org.postgresql.Driver)

URL

jdbc:postgresql://database-1.cjpk2s29ue0.ap-south-1.rds.amazonaws.com:5432/PROD

Username

puser

Password

ppassword

Autocommit

☒

Fetch size

Timeout

s

SSH

Extended Properties

Connect

4. Load Data to create all tables in Postgre

02.Postgre_DDL.sql

schemaname	tablename	tableowner	tablespace	hasindexes	hasrules	hastriggers	rowsecurity
public	country	puser		true	false	false	false
public	city	puser		true	false	false	false
public	address	puser		true	false	false	false
public	card_type	puser		true	false	false	false
public	card	puser		true	false	false	false
public	cc_debit	puser		false	false	false	false
public	tx_type	puser		true	false	false	false
public	cc_paid	puser		false	false	false	false

5. Initial Data load using

03.Initial_Dataload.sql

Check table structure and sample data

6. Create EMR (5.33) with Below Software's

- Hadoop
- Spark
- Hbase
- Zookeeper
- Hive
- Hue
- Phoenix

7. Copy All dependencies on EMR

/home/hadoop/dep/*

8. Create Hbase table using phoenix Script

04.Pheonix_Hbase_DDL_CCDA.txt

```
CREATE TABLE COUNTRY_HB (cn_id VARCHAR(10) PRIMARY KEY,cn_name VARCHAR(45));
CREATE TABLE CITY_HB (ct_id VARCHAR(10) PRIMARY KEY,ct_name VARCHAR(45),cn_id VARCHAR(10));
CREATE TABLE CARD_TYPE_HB (c_type VARCHAR(10) PRIMARY KEY,max_limit INTEGER);
CREATE TABLE TX_TYPE_HB (tx_type_id VARCHAR(10) PRIMARY KEY,tx_type_desc VARCHAR(45));
```

Note : To run pyspark job on EMR cluster you can use below command :

```
/usr/bin/spark-submit --jars /home/hadoop/dep/postgresql-
42.2.14.jar,/home/hadoop/dep/phoenix-4.14.3-HBase-1.4-client.jar,/home/hadoop/dep/phoenix-
spark-4.14.3-HBase-1.4.jar --master yarn --deploy-mode client --driver-memory 3g --executor-memory
2g --num-executors 1 --executor-cores 1 /home/hadoop/ 05.CCDA_RDS_to_HBASE_Ref_import.py
```

9. Import Ref data from RDBMS to Hbase using (Change the DB details in Script)

05.CCDA_RDS_to_HBASE_Ref_import.py

```
#dbtable="COUNTRY"

print("Data import for reference table is started....")
#Country Table
df_cn=spark.read.format("jdbc").option("url",host).option("user",user).option("password",pwd).option("driver"
,driver).option("dbtable","COUNTRY").load()
df_cn.write.format("org.apache.phoenix.spark").option("table","COUNTRY_HB").option("zkUrl","localhost:2181").
mode('overwrite').save()
print("Country Table Imported successfully")
```

Sample Data :

```
0: jdbc:phoenix:localhost> select * from country_hb limit 10;
```

CN_ID	CN_NAME
CN_1	Afghanistan
CN_10	Azerbaijan
CN_100	Ukraine
CN_101	United Arab Emirates
CN_102	United Kingdom
CN_103	United States
CN_104	Venezuela
CN_105	Vietnam
CN_106	Virgin Islands, U.S.
CN_107	Yemen

10. Import Trans data from RDBMS to Hbase using (Change the DB details in Script)

06.RDS_to_HBASE_Hist_Trans_import.py

```
print("Historical Data import for transactional table is started....")

#CARD Table
df_sb=spark.read.format("jdbc").option("url",host).option("user",user).option("password",pwd).option("driver",driver).option("dbtable","CARD").load()
df_sb.write.format("org.apache.phoenix.spark").option("table","CARD_HB").option("zkUrl","localhost:2181").mode('overwrite').save()
print("CARD Table Imported successfully")
```

Sample Data :

C_NUMBER	FULL_NAME	MOB	EMAIL	ISSUE_DATE	UPDATE_DATE	EXP_DATE	BIL
111111	MARY SMITH	54321	MARY.SMITH@sakilacustomer.org	2021-01-01 00:00:00.000	2021-01-01 00:00:00.000	2026-01-01 00:00:00.000	1
111112	PATRICIA JOHNSON	54322	PATRICIA.JOHNSON@sakilacustomer.org	2021-01-01 00:00:00.000	2021-01-01 00:00:00.000	2026-01-01 00:00:00.000	1
111113	LINDA WILLIAMS	54323	LINDA.WILLIAMS@sakilacustomer.org	2021-01-01 00:00:00.000	2021-01-01 00:00:00.000	2026-01-01 00:00:00.000	1
111114	BARBARA JONES	54324	BARBARA.JONES@sakilacustomer.org	2021-01-01 00:00:00.000	2021-01-01 00:00:00.000	2026-01-01 00:00:00.000	1
111115	ELIZABETH BROWN	54325	ELIZABETH.BROWN@sakilacustomer.org	2021-01-01 00:00:00.000	2021-01-01 00:00:00.000	2026-01-01 00:00:00.000	1
111116	JENNIFER DAVIS	54326	JENNIFER.DAVIS@sakilacustomer.org	2021-01-01 00:00:00.000	2021-01-01 00:00:00.000	2026-01-01 00:00:00.000	1
111117	MARIA MILLER	54327	MARIA.MILLER@sakilacustomer.org	2021-01-01 00:00:00.000	2021-01-01 00:00:00.000	2026-01-01 00:00:00.000	1
111118	SUSAN WILSON	54328	SUSAN.WILSON@sakilacustomer.org	2021-01-01 00:00:00.000	2021-01-01 00:00:00.000	2026-01-01 00:00:00.000	1
111119	MARGARET MOORE	54329	MARGARET.MOORE@sakilacustomer.org	2021-01-01 00:00:00.000	2021-01-01 00:00:00.000	2026-01-01 00:00:00.000	1
111120	DOROTHY TAYLOR	54330	DOROTHY.TAYLOR@sakilacustomer.org	2021-01-01 00:00:00.000	2021-01-01 00:00:00.000	2026-01-01 00:00:00.000	1

11. Create Hive tables deployment to transformed data using

07.CCDA_hive_ddl.hql

```
create table card_details_staging (CardNumber integer,CardType string,full_name string,contactnumber integer,emailId string,address string,city string,country string,issuedate timestamp,update_date timestamp,BillingDate integer,CardLimit integer,Active_flag Char(1))
Row format delimited
fields terminated by ',';
```

```
create table card_details (CardNumber integer,CardType string,full_name string,contactnumber integer,emailId string,address string,city string,country string,issuedate timestamp,update_date timestamp,BillingDate integer,CardLimit integer,Active_flag Char(1))
Row format delimited
fields terminated by ',';
```

Tables in Hive :

```
hive> show tables;
OK
card_details
card_details_staging
credit_details
credit_details_staging
debit_details
debit_details_staging
Time taken: 0.168 seconds, Fetched: 6 row(s)
```

12. To process Card details context using below script

08.CCDA_Hb_to_hive_card_details.py

```
#DSL
test=df_cd_tmp.join(df_ad_tmp,"add_id", how="left").join(df_ct_tmp,"ct_id", how="left").join(df_cn_tmp,
"cn_id", how="left").drop("ADD_ID").drop("CT_ID").drop("CN_ID")

#Give alias

res=test.selectExpr("C_NUMBER as CardNumber", "C_TYPE as CardType","FULL_NAME as CardHolderName","MOB as
ContactNumber","EMAIL as emailId","STREET as address","CT_NAME as city","CN_NAME as country","ISSUE_DATE as
issudate","UPDATE_DATE as update_date","BILLING_DATE as BillingDate","C_LIMIT as CardLimit","ACT_FLAG as
Active_flag")

res.show(5)
```

Sample Data :

```
hive> select * from card_details_staging limit 5;
OK
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
111421 Gold PAUL TROUT 54631 PAUL.TROUT@sakilacustomer.org 746 Joliet Lane Kursk Russian Federation 2021-01-01 00:00:00 2021-01-01 00:00:00 1
7 90000 A
112020 Business EASTER BEN 55230 EASTER.BEN@customer.org 746 Joliet Lane Kursk Russian Federation 2021-01-01 00:00:00 2021-01-01 00:00:00 1
7 400000 A
111330 Gold CHARLENE ALVAREZ 54540 CHARLENE.ALVAREZ@sakilacustomer.org 1842 Luzinia Boulevard Zanzibar Tanzania 2021-01-01 00:00:00 2021-01-01 00:00:00 2
021-01-01 00:00:00 17 90000 A
111696 Titanium KIRK STCLAIR 54906 KIRK.STCLAIR@sakilacustomer.org 1923 Stara Zagora Lane Tsaotun Taiwan 2021-01-01 00:00:00 2021-01-01 00:00:00 1
7 100000 A
111929 Business CARBONE CHAD 55139 CARBONE.CHAD@customer.org 1842 Luzinia Boulevard Zanzibar Tanzania 2021-01-01 00:00:00 2021-01-01 00:00:00 1
01 00:00:00 17 400000 A
Time taken: 1.537 seconds, Fetched: 5 row(s)
hive>
```

13. To process debit and Credit details context use below script

09.CCDA_Hb_to_hive_debit_details.py

```
hive> select * from debit_details_staging;
OK
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
111111 TD000001      2021-02-01 00:00:00      1000      Online      Success
111111 TD000002      2021-02-01 00:00:00      500       Food        Success
111111 TD000003      2021-02-02 00:00:00      2000      Travel      Success
111112 TD000004      2021-02-01 00:00:00      1000      Online      Success
111112 TD000005      2021-02-02 00:00:00      500       Food        Success
111112 TD000006      2021-02-02 00:00:00      1500      Travel      Success
111112 TD000007      2021-02-03 00:00:00      2000      Fuel        Success
Time taken: 1.368 seconds, Fetched: 7 row(s)
hive>
```

10.CCDA_Hb_to_hive_credit_details.py

```
hive> select * from credit_details_staging limit 5;
OK
111112 RX_0000002      2021-02-10 00:00:00      2000      success Bank Application
111111 RX_0000001      2021-02-10 00:00:00      2000      success Netbanking
Time taken: 0.12 seconds, Fetched: 2 row(s)
hive>
```

14. Remove duplicate data using

10.dedup_Compaction.py

```
print("Card details de-duplication and compaction started")
cd_dff=spark.sql("select * from prod.card_details")
cd_dfs=spark.sql("select * from prod.card_details_staging")

un = cd_dff.union(cd_dfs)

res=un.withColumn("new_upd",F.when(un.update_date.isNull(),F.to_timestamp(F.lit("1970-01-01 00:00:00 "),
format="yyyy-MM-dd HH:mm:ss")).otherwise(un.update_date)).drop("update_date")

f=res.select("*").withColumnRenamed("new_upd","update_date")

res1=f.withColumn("rn", F.row_number().over(Window.partitionBy("CardNumber").orderBy(desc("update_date"))))

res2=res1.filter(res1.rn == 1).drop("rn")

res2.show(30)
```

15. To Process delta use below script

Copy data into Hbase at /data location

Run this script to process new delta

11.Dedup_CCDA.py

Delta :

Address :

```
ADD_ID,STREET,CT_ID
AD_66,1717 Guadalajara Lane,CT_441
```

Card :

```
C_NUMBER,FULL_NAME,MOB,EMAIL,ISSUE_DATE,UPDATE_DATE,EXP_DATE,BILLING_DATE,C_LIMIT,ACT_FLAG,ADD_ID,C_TYPE
111112,PATRICIA
JOHNSON,54322,PATRICIA.JOHNSON@sakilacustomer.org,2021-01-01,2021-01-01,2026-01-01,1,120000,A,AD_66,Gold
999999,Rahul
Shinde,99999,rahul.shinde@sakilacustomer.org,20-21-08-01,20-21-08-01,2026-01-01,1,120000,A,AD_66,Visa
```

Debit :

```
C_NUMBER, TX_ID, TX_DATE, AMT_SPEND, CATEGORY, D_STATUS
111112, TD000008, 2021-08-01, 1000, Shopping, success
```

Credit:

```
C_NUMBER, TX_ID, TX_DATE, AMT_PAID, C_STATUS, TX_TYPE_ID
111112, RX_0000003, 2021-08-2, 2000, success, APP
```

Script to process delta :

```
card_path="/ccda_data/card/*"
debit_path="/ccda_data/debit/*"
credit_path="/ccda_data/credit/*"
caddress_path="/ccda_data/caddress/*"

#Storing data into Hbase
df_card=spark.read.format("csv").option("header","true").option("inferSchema","true").load(card_path)
df_card.write.format("org.apache.phoenix.spark").option("table","CARD_HB").option("zookeeper","localhost:2181").
mode('overwrite').save()
```

```

print("Processing Context for Card Details")

df_result_card=spark.sql("""SELECT
    C.c_number as CardNumber,
    C.c_type as CardType,
    C.c_full_name as CardHolderName,
    C.mob as contactnumber,
    C.email as emailId,
    Ad.street as address,
    ct.ct_name as city,
    cn.cn_name as country,
    C.issue_date as issuedate,
    C.update_date as update_date,
    C.billing_date as BillingDate,
    C.c_limit as CardLimit,
    C.act_flag as Active_flag
from credit_tmpl C
LEFT JOIN Address_tmpl Ad on C.add_id=Ad.add_id
Left Join City_tmpl ct on Ad.ct_id=ct.ct_id
Left join Country_tmpl cn on ct.cn_id=cn.cn_id

""")

df_result_card.show(10)

```

```

hive> select * from card_details_staging;
OK
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
111112 Gold PATRICIA JOHNSON 54322 PATRICIA.JOHNSON@sakilacustomer.org 1717 Guadalajara Lane Saint Louis United States 2021-01-01 00:00:00 2
021-01-01 00:00:00 1 120000 A
999999 Visa Rahul Shinde 99999 rahul.shinde@sakilacustomer.org 1717 Guadalajara Lane Saint Louis United States 2021-08-01 00:00:00 2021-08-01 00:00
:00 1 120000 A
Time taken: 1.384 seconds, Fetched: 2 row(s)
hive>

```

Note : You have successfully Processed Historical and Delta Data

Phase 2 :

Extraction :

```
bill_report=spark.sql("""
SELECT
CD.CardNumber as cardNumber,
CD.CardType as CardType,
CD.full_name as full_name,
CD.contactnumber as mobile,
CD.emailId as email,
concat(CD.address,'_',CD.city,'_',CD.country) as cust_address,
CD.BillingDate as billdate,
CD.CardLimit as CardLimit,
CD.cardLimit-D.totalDebitAmount+C.totalCreditAmount as outstandingLimit,
D.totalDebitAmount-C.totalCreditAmount as outstandingAmount
from prod.card_details CD
LEFT JOIN
(select cardNumber,sum(DebitAmount) as totalDebitAmount from prod.debit_details
where UPPER(debitTxStatus)='SUCCESS' group by cardNumber)D on
CD.cardNumber=D.CardNumber
LEFT JOIN
(select cardNumber,sum(CreditAmount) as totalCreditAmount from prod.credit_details
where UPPER(CreditTxStatus)='SUCCESS' group by cardNumber) C ON
CD.cardNumber=C.CardNumber
where CD.Active_flag='A' and CD.cardLimit-D.totalDebitAmount+C.totalCreditAmount <>
CD.cardLimit
""")
```

Sample Data :

Now customer can extract report from S3 and can use for their further use like visualization etc.

cardNumber	CardType	full_name	mobile	email	cust_address	billdate	CardLimit	outstandingLimit	outstandingAmount
111111	Platinum	MARY SMITH	54321	MARY.SMITH@sakilacustomer.org	1913 Hanoi Way_Sasebo_Japan	1	45000	43500	1500
111112	Visa	PATRICIA JOHNSON	54322	PATRICIA.JOHNSON@sakilacustomer.org	1121 Lida Avenue_San Bernardino_United States	1	120000	118000	2000

Job Monitoring on Resource Manager :

<http://ec2-13-235-62-102.ap-south-1.compute.amazonaws.com:8088/cluster>

[All Applications :](#)

New Tab x FoxyProxy options x All Applications x +

Not secure | ec2-13-235-62-102.ap-south-1.compute.amazonaws.com:8088/cluster

Apps | Scala_Spark | LinkedIn | MyNokari7794 | MyNokari520 | sagaringale7794@g... | sagaringale520@g... | Apache Hadoop on... | Quara | Other bookmarks | Reading list

hadoop

All Applications

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Used Resources	Total Resources	Reserved Resources
10	0	0	10	0	<memory:0, vCores:0>	<memory:5144, vCores:4>	<memory:0, vCores:0>

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
1	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Capacity Scheduler	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[<memory:memory.mb.default.unitMM.type=COUNTABLE>, <memory:vcores.default.unit=type=COUNTABLE>]	<memory:0, vCores:1>	<memory:5144, vCores:4>	0

Showing 1 to 10 of 10 entries

ID	User	Name	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinalTime	State	FinalStatus	Running Containers	Allocated CPU vCores	Allocated Memory MB	Allocated GPUs	Reserved CPU vCores	Reserved Memory MB	Reserved GPUs	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes
application_1628614480454_0010	hadoop	Total_revenue_report	SPARK	default	0	Tue Aug 10 23:24:44 +0550	Tue Aug 10 23:24:44 +0550	Tue Aug 10 23:25:03	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	100%	History	0
application_1628614480454_0009	hadoop	HIVE-0869446a-b79-435b-886a-903a57c50926	TEZ	default	0	Tue Aug 10 23:17:49 +0550	Tue Aug 10 23:17:50 +0550	Tue Aug 10 23:23:28	FINISHED	ENDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	100%	History	0
application_1628614480454_0008	hadoop	Tel_dekup_Compaction	SPARK	default	0	Tue Aug 10 23:16:59 +0550	Tue Aug 10 23:16:59 +0550	Tue Aug 10 23:18:32	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	100%	History	0
application_1628614480454_0007	hadoop	Tel_dekup_Compaction	SPARK	default	0	Tue Aug 10 23:16:02 +0550	Tue Aug 10 23:16:02 +0550	Tue Aug 10 23:18:32	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	100%	History	0
application_1628614480454_0006	hadoop	hdfsive_sub_data1s	SPARK	default	0	Tue Aug 10 23:03:33 +0550	Tue Aug 10 23:03:33 +0550	Tue Aug 10 23:04:03	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	100%	History	0
application_1628614480454_0005	hadoop	HIVE-0869446a-b79-435b-886a-903a57c50926	TEZ	default	0	Tue Aug 10 23:02:43 +0550	Tue Aug 10 23:02:43 +0550	Tue Aug 10 23:12:23	FINISHED	ENDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	100%	History	0
application_1628614480454_0004	hadoop	hdfsive	SPARK	default	0	Tue Aug 10 22:55:28 +0550	Tue Aug 10 22:55:28 +0550	Tue Aug 10 22:55:56	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	100%	History	0
application_1628614480454_0003	hadoop	hdfsive_sub_data1s	SPARK	default	0	Tue Aug 10 22:53:32 +0550	Tue Aug 10 22:53:32 +0550	Tue Aug 10 22:53:53	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	100%	History	0
application_1628614480454_0002	hadoop	HIVE-0869446a-b79-435b-886a-903a57c50926	TEZ	default	0	Tue Aug 10 22:44:57 +0550	Tue Aug 10 22:44:57 +0550	Tue Aug 10 22:55:06	FINISHED	ENDED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	100%	History	0
application_1628614480454_0001	hadoop	HIVE-6162b2b1-8733-417d-b1ad-4b3847c10df9	TEZ	default	0	Tue Aug 10 22:44:43 +0550	Tue Aug 10 22:44:44 +0550	Tue Aug 10 22:44:45	KILLED	KILLED	N/A	N/A	N/A	N/A	N/A	N/A	N/A	0.0	0.0	100%	History	0

Showing 1 to 10 of 10 entries

Application Logs :

hadoop

Application application_1628614480454_0010

Logged in as: dr.whe

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler
- Tools

Application Overview

User: hadoop

Name: Total_revenue_report

Application Type: SPARK

Application Tags:

Application Priority: 0 (Higher Integer value indicates higher priority)

YarnApplicationState: FINISHED

Queue: default

FinalStatus Reported by AM: SUCCEEDED

Started: Tue Aug 10 17:54:44 +0000 2021

Launched: Tue Aug 10 17:54:44 +0000 2021

Finished: Tue Aug 10 17:55:03 +0000 2021

Elapsed: 18sec

Tracking URL: History

Log Aggregation Status: SUCCEEDED

Application Timeout (Remaining Time): Unlimited

Diagnostics:

Unmanaged Application: false

Application Node Label expression: <Not set>

AM container Node Label expression: CORE

Application Metrics

Total Resource Preempted: <memory:0, vCores:0>

Total Number of Non-AM Containers Preempted: 0

Total Number of AM Containers Preempted: 0

Resource Preempted from Current Attempt: <memory:0, vCores:0>

Number of Non-AM Containers Preempted from Current Attempt: 0

Aggregate Resource Allocation: 86520 MB-seconds, 46 vcore-seconds

Aggregate Preempted Resource Allocation: 0 MB-seconds, 0 vcore-seconds

Showing 20 entries

Attempt ID	Started	Node	Logs	Nodes blacklisted by the app	Nodes blacklisted by the system
appattempt_1628614480454_0010_000001	Tue Aug 10 23:24:44 +0550 2021	http://ec-172-31-17-143.ap-south-1.compute.internal:8042	Logs	0	0

Showing 1 to 1 of 1 entries

First Previous 1 Next Last

Check Log on Terminal :

```

EEEEEEEEEEEEEEEEEEEE MMMMMMM      MMMMMMM RRRRRRRRRRRRRR
E::::::::::::::::::::E M::::::::M      M::::::::M R:::::::::R
EE::::::::EEEEEEEE::E M::::::::M      M::::::::M R::::RRRRR:::R
  E:::E      EEEE M::::::::M      M::::::::M RR:::R      R:::R
  E:::E      M:::M M:::M M:::M M:::M M:::M R:::R      R:::R
  E:::EEEEEEEE M:::M M:::M M:::M M:::M R::RRRRR:::R
  E:::EEEEEEEE M:::M M:::M M:::M R::RRRRR:::R
  E:::E      M:::M M:::M M:::M R::R      R:::R
  E:::E      EEEE M:::M      M M:::M M:::M R::R      R:::R
EE::::::::EEEEEEEE::E M:::M      M:::M R:::R      R:::R
E::::::::::::::::::::E M:::M      M:::M RR:::R      R:::R
EEEEEEEEEEEEEEEEEEEE MMMMMMM      MMMMMMM RRRRRR      RRRRRR

```

```
[hadoop@ip-172-31-17-143 ~]$ yarn logs -applicationId application_1628614480454_0010
```

Error Occurred :

```

py4j.protocol.Py4JJavaError: An error occurred while calling o75.load.
: java.lang.ClassNotFoundException: org.postgresql.Driver
    at java.net.URLClassLoader.findClass(URLClassLoader.java:382)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:418)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:351)
    at org.apache.spark.sql.execution.datasources.jdbc.DriverRegistry$.register(DriverRegistry.scala:45)
    at org.apache.spark.sql.execution.datasources.jdbc.JDBCOptions$$anonfun$$apply(JDBCOptions.scala:99)
    at org.apache.spark.sql.execution.datasources.jdbc.JDBCOptions$$anonfun$$apply(JDBCOptions.scala:99)
    at scala.Option.foreach(Option.scala:257)
    at org.apache.spark.sql.execution.datasources.jdbc.JDBCOptions.<init>(JDBCOptions.scala:99)
    at org.apache.spark.sql.execution.datasources.jdbc.JDBCOptions.<init>(JDBCOptions.scala:35)
    at org.apache.spark.sql.execution.datasources.jdbc.JdbcRelationProvider.createRelation(JdbcRelationProvider.scala:32)
    at org.apache.spark.sql.execution.datasources.DataSource.resolveRelation(DataSource.scala:332)
    at org.apache.spark.sql.DataFrameReader.loadV1Source(DataFrameReader.scala:242)
    at org.apache.spark.sql.DataFrameReader.load(DataFrameReader.scala:230)
    at org.apache.spark.sql.DataFrameReader.load(DataFrameReader.scala:186)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
    at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:357)
    at py4j.Gateway.invoke(Gateway.java:282)
    at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:132)
    at py4j.commands.CallCommand.execute(CallCommand.java:79)
    at py4j.GatewayConnection.run(GatewayConnection.java:238)
    at java.lang.Thread.run(Thread.java:749)

21/08/10 18:11:00 INFO SparkContext: Invoking stop() from shutdown hook
21/08/10 18:11:00 INFO SparkUI: Stopped Spark web UI at http://ip-172-31-17-143.ap-south-1.compute.internal:4040
21/08/10 18:11:00 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
21/08/10 18:11:00 INFO MemoryStore: MemoryStore cleared
21/08/10 18:11:00 INFO BlockManager: BlockManager stopped
21/08/10 18:11:00 INFO BlockManagerMaster: BlockManagerMaster stopped
21/08/10 18:11:00 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
21/08/10 18:11:00 INFO SparkContext: Successfully stopped SparkContext
21/08/10 18:11:00 INFO ShutdownHookManager: Shutdown hook called
21/08/10 18:11:00 INFO ShutdownHookManager: Deleting directory /mnt/tmp/spark-6faa3d5f-f74a-47fb-bed2-d9fc449be06a
21/08/10 18:11:00 INFO ShutdownHookManager: Deleting directory /mnt/tmp/spark-cead1d44-cl91-4634-806b-1c18639cec2c/pyspark-a61096cf-394d-4345-a23d-d8c8255eb710
21/08/10 18:11:00 INFO ShutdownHookManager: Deleting directory /mnt/tmp/spark-cead1d44-cl91-4634-806b-1c18639cec2c
[hadoop@ip-172-31-17-143 ~]$

```

UI Port number :

<https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-web-interfaces.html>

Name of interface	URI
Flink history server (EMR version 5.33 and later)	http://master-public-dns-name:8082/
Ganglia	http://master-public-dns-name/ganglia/

Name of interface	URI
Hadoop HDFS NameNode (EMR version pre-6.x)	https:// <i>master-public-dns-name</i> :50470/
Hadoop HDFS NameNode	http:// <i>master-public-dns-name</i> :50070/
Hadoop HDFS DataNode	http:// <i>coretask-public-dns-name</i> :50075/
Hadoop HDFS NameNode (EMR version 6.x)	https:// <i>master-public-dns-name</i> :9871/
Hadoop HDFS DataNode (EMR version pre-6.x)	https:// <i>coretask-public-dns-name</i> :50475/
Hadoop HDFS DataNode (EMR version 6.x)	https:// <i>coretask-public-dns-name</i> :9865/
HBase	http:// <i>master-public-dns-name</i> :16010/
Hue	http:// <i>master-public-dns-name</i> :8888/
JupyterHub	https:// <i>master-public-dns-name</i> :9443/
Livy	http:// <i>master-public-dns-name</i> :8998/
Spark HistoryServer	http:// <i>master-public-dns-name</i> :18080/
Tez	http:// <i>master-public-dns-name</i> :8080/tez-ui
YARN NodeManager	http:// <i>coretask-public-dns-name</i> :8042/
YARN ResourceManager	http:// <i>master-public-dns-name</i> :8088/
Zeppelin	http:// <i>master-public-dns-name</i> :8890/